

Official SWE-rebench 111: who we ran, and what we see so far

Snapshot 31 August 2026, 10:44 UTC. Run still in progress. Numbers will move. Do not treat this as a final board result.

Status. All 111 tasks have at least one try. Genesis is about 73% through its 555 trials (111 tasks × 5 tries). King and teacher lag. Extra tries will change $\text{pass}@k$ more than the mean.

Why this run exists

Affine (Bittensor subnet 120) ranks miners with a teacher-anchored score called **min(R,G)**. That score never touches SWE-rebench tasks. It asks: does the miner’s thought help the teacher predict the teacher’s next action, and does that thought look like a real teacher thought?

SWE-rebench is a *separate* coding exam, published at swe-rebench.com. We ran it so we can say, in public, how the models on our throne compare to the official board — not to our older 860-task Python dump, which is a different exam.

What we ran

The official July 2026 window: **111 tasks** (15 May–1 July 2026), 65 repositories, five languages (Go, Java, Python, Rust, TypeScript). Each model gets **5 independent tries** per task. We report the **mean resolve rate ± SEM** (the board headline) and **pass@k** (share of tasks solved at least once in k tries).

Knob	Official board	This run
Task set	111-task July 2026 window	Same Harbor dataset (swe-rebench-07-2026)
Context / step cap	128k context, 300 steps	Same (vLLM max-model-len 131072)
Sampling	Vendor defaults (Qwen: temp 0.6, top_p 0.95, top_k 20)	Same
Agent text	Published ReAct prompt with `` command fences	Same published text
Editor tools	Nebius open/edit/submit wrappers (unpublished)	Ordinary bash. No official editor binary.
Scoring	5 tries; mean ± SEM; pass@5	Same plan. Snapshot below is still filling tries.

Honest gap: we cannot run Nebius’s private editor. Calibration is the check. Official untouched Qwen3.6-35B-A3B on this exam is **24.7% ± 0.79%** mean, **43.2%** pass@5. If our genesis lands near that, the protocol is close enough to publish. If it stays near 17%, we are still on a weaker scaffold.

Who the models are

Role	Weights	What it means on SN120
Genesis	Qwen/Qwen3.6-35B-A3B revision 995ad96e	Untouched base model. Seed king of the min(R,G) era. Unpaid until a miner crowns. Official SWE-rebench number for this exam: 24.7%. The calibration stick.
King	vera6/affine-5g4yy75zuz-q1 revision dbac4185	First miner that beat genesis on live min(R,G). A genesis-family fine-tune. Winning our duel does not automatically mean it writes better patches.
Teacher	Qwen/Qwen3.8-27B	The model that scores miners. It writes the reference thoughts and actions. It is not on the SN120 board. We run it to see if the judge is a stronger coding agent than the models it ranks.

Results so far (31 Aug 2026, 10:44 UTC)

Role	Trials done	Mean resolve	Pass@k so far	Vs official genesis
Genesis	403 / 555 (~4 tries/task)	17.6% $\pm$ 2.9%	pass@4 30.6%	Official: 24.7% $\pm$ 0.79% pass@5 43.2% We are ~7 points low
King	310 / 555 (~3 tries/task)	12.5% $\pm$ 2.4%	pass@3 24.3%	Behind genesis on this exam
Teacher	135 / 555 (~2 tries/task)	44.1% $\pm$ 4.6%	pass@2 47.7%	Far ahead of both throne models

Mean resolve is the average, over the 111 tasks, of (successful tries $\div$ tries so far). That is the official headline. **Pass@k** is the share of tasks solved at least once in k tries. Our k is not yet 5 for every model. Extra tries usually raise pass@k. They do not automatically raise the mean toward 24.7%.

How to read this

1. The king is **not** beating genesis on the official exam. The earlier 860-task Python dump (king 23%, teacher 22%, genesis 15%) was a different set, one try, a shorter agent, and Python only. That ranking is not repeating here.
2. The teacher 27B is a much stronger coding agent than the 35B-class models on the throne. That fits Qwen's own agentic-coding numbers. It does *not* mean we should put the teacher on the SN120 board. The board is pinned to the genesis family on purpose, so the first teacher upload cannot take the throne by construction.
3. Do not put 17.6% next to Fable 5 / Grok / Opus (~62–65%) and call it a fair fight. Those numbers are on this 111-task exam with Nebius's own agent. Our genesis is also below their published 24.7% for the same weights. Until that gap closes or we state it, these are **internal comparable ranks**, not board-ready scores.

What we will not claim

We will not claim coding isomorphism (higher min(R,G) $\Rightarrow$ higher SWE-rebench). The live SN120 panel already inverted that story once (RT-7). This run is a measurement, not a proof of the meter.

We will not claim this snapshot is final. Teacher still has most of its 555 trials left. Timeouts (50-minute agent cap) hit the teacher hardest. Remaining tries can move pass@k by several points.

Affine SN120 · unofficial measurement · Harbor + mini-swe-agent on the July 2026 official window · not a Nebius submission · weight_version_key unchanged (10). Contact: affine.io